

Performance Calibrations at Tech Companies: Part 1

The reasons for performance calibrations, budgets, bucketing and stack ranking approaches.



GERGELY OROSZ

DEC 6, 2022 • PAID



82



5



Share



Performance reviews and promotions are coming soon at many tech companies, within a month or two. We've already covered [Preparing for performance reviews ahead of time](#) and [Preparing for promotions ahead of time](#) in previous articles. But in this issue, we dive into a topic highly relevant for engineering managers, which is also helpful to understand for individual contributors: performance calibrations.

'Performance calibration' refers to the meeting when managers gather and decide which rating – and bonuses and pay rises – to assign to people. So, what really happens inside this crucial meeting?

For this question, I'm bringing my own experience of being in many calibration sessions at Uber, as well as input from half a dozen engineering directors, CTOs and senior engineering managers at tech companies.

It's hard to appreciate how performance calibrations work without understanding their purpose and the constraints by which they operate. Today, we cover these topics:

1. **"Why?"** Why do performance calibrations exist at most companies? Which places operate without them?
2. **Companies operating without formal performance processes.** For how long can this setup work at places with no standard performance processes, and therefore no calibrations?

3. **Budgets.** Why budgets are important to understand, and how they relate to performance calibrations and budget distribution strategies.
4. **Performance calibrations.** Goals of the calibration process and the various types of calibration across the tech industry.
5. **Bucketing.** How does this work, and what are the typical performance buckets at tech companies?
6. **Stack ranking approaches.** How this works and its relation to calibration.

Additionally, [Part 2 also covers:](#)

- **A “typical” Big Tech calibration.** What happens? Why do calibrations take weeks to complete?
- **Preparing ahead of time.** As a manager – or a manager of managers – how can you prepare for calibration?
- **Inside a calibration.** Politics, deadlocks, strategies, allies. What are the typical situations which arise, and how can you respond to them as a manager?
- **Biases during calibration.** Common biases at play during calibration, how to spot and counter them as a manager in the room. Also, ways you could help promote fairness by pushing back against biases.
- **Advice for managers.** How can you set up your direct reports to get better feedback from others during the calibration? Also, what should you share with your directs and what should you keep private?
- **Advice for individual contributors (ICs)** in preparing for calibrations, ahead of time. It turns out, you can tip the scale in your favor by helping your manager, if only by a limited amount.

1. The “why” of performance management

If you’ve worked at a larger company, you’ll know that formal performance reviews – together with a promotion process – are usually parts of how these workplaces operate. But let’s take a step back and ask the question: why? Why is there a

performance process at pretty much every company? To find out, let's do a thought experiment about a company with no performance management process.

Let's assume this business starts as a small startup with a few cofounders who are friends, who work hard, and then hire the first few employees. Those colleagues also work hard and the company thrives.

However, as the company grows employees start to ask questions, like:

- Why does Employee X or Y make more money than me, when we do roughly the same work?
- How big a raise will I get? Will everyone get the same raise? If so, how is that fair when I worked longer and harder than some others? If the raises are not equal, who decides and based on what, who gets the bigger raise?
- Why does a new hire get paid a lot more than me, when we do roughly the same job?

These questions are common and are about two things:

- **Expectations.** What's expected from me, versus what is expected of others?
- **Rewards and compensation.** What can I expect if I do an okay job, a great job, or if I put my heart and soul into this role?

In our thought experiment, a few things happen when these questions go unaddressed:

- People grumble more and more, usually that:
 - Things are opaque as everyone's compensation is different, and compensation levels boil down to an individual's negotiating skill.
 - Things are unfair when everyone is paid the same and raises are identical, as those who do more or better work are rewarded the same as those who slack off.
- Employees who feel they work harder than others but don't get rewarded for it, may often leave. These people are often the top performers.

- Those who don't pull their weight and do far less work than others, often stay with the company. Such people make it harder to get things done, pull down team morale, and drive some higher performing colleagues to eventually seek opportunities at workplaces where they don't have to put up with low performing co-workers and are better rewarded.

Without performance management in place, a company will find it close to impossible to identify who's a top or low performer, or at least will have trouble doing so fairly:

- **The low performers.** Those doing less or poorer quality work than expected. They might be slacking off or coasting, or they might lack the skills needed for their job.
- **The top performers.** People who stand out for their contributions. Among software engineers, these are the ones who tend to ship things faster, get things done with fewer bugs and regressions, take actions which reduce the effort of shipping, and lift up colleagues in their team and around them.
- **Those who are ready to be promoted to the next level.** Above a certain size, companies introduce levels for roles like software engineer and engineering manager. The higher the level, the greater the compensation and the bigger the expectations are. Making it clear where people are relative to the next level, becomes an important way to help them grow, to motivate them, and to keep them grounded – especially early during careers. Read more about promotions in the issue [Preparing for promotions ahead of time](#).

Low performers are folks who need to hear this feedback and to improve. Top performers also should get feedback and support for career opportunities and growth.

A performance management process helps identify low and top performers. Conveniently, it can also resolve compensation questions by providing a predictable, systematic answer to the question: "why has person X got a higher/lower raise."

2. Companies operating without formal performance processes

There are some companies that operate with no performance management process. I talked with a software engineer working at one such business, who shared:

"Not having a performance management process is both good and bad. On the good part: there's a lot less stress.

On the bad part: underperformers stick around, barely doing any work. It is unclear why people are promoted. People who are poor leaders or managers stay at the company. Basically: those who do poor work have a stable job here."

Common types of companies without a performance management process, include:

- **Early-stage startups.** Typically less than two years old, where putting a process in place makes little sense as yet.
- **Digital agencies.** Digital agencies or developer agencies contract with companies to do development work. They frequently hire less experienced software engineers, sometimes train them, and will contract these people out for projects, or on an hourly basis. It's common these agencies have no performance management process, especially the smaller ones. Larger businesses such as Accenture, EPAM or Thoughtworks, will have both leveling and performance processes.
- **Older non-tech companies** in several countries. I've talked with software engineers, including in Germany, who note that many companies they'd describe as "ancient" and not in tech, work with no performance management processes.

Scaleups are among companies that do performance management "on the fly" without a formal, annual process.

A fintech scaleup founded in 2007 and acquired for more than \$5B just over 10 years later, had no annual reviews before its acquisition. Managers would have weekly catch ups with reports to address issues they might have seen with performance. Those doing well would be promoted relatively quickly, while those not doing well exited the business.

A former software engineer at this scaleup shared that management cared about unusual metrics like the number of comments on pull requests and code statistics. This person speculated that performance reviews were overlooked for the same reason as sick leave was not tracked: leadership didn't think it was worth the time.

At this company, the artificial coding metrics that managers and directors occasionally looked at were evidently the most important measures of success. This software engineer observed:

"The worst engineers were promoted to engineering manager and director levels. On paper, they had lots of pull requests, even though they often went straight to production and broke things. As the business grew, we ended up with a layer of leadership who were great while we were a startup, but terrible in a regulated environment."

Let's be honest: a place with no performance management system can easily become a hotbed of favoritism, nepotism and even cronyism. Low performers will not be managed out, and only those lacking the "right" connections will get the boot.

So, performance management is a "necessary evil" for any healthy organization.

It's not that managers or employees enjoy conducting the process up to four times a year, depending on the company. It's just that without a process to gather and share feedback in a formal way, and the ability to be clear on whether or not people are meeting expectations, you'll end up with cronyism or favoritism in a place where underperformers are welcome to stay – or a corrosive mix of them all.

Performance management processes are in place across Big Tech, and are all fairly similar. We previously covered in-depth [how it works at Facebook](#), and [performance management at Amazon](#).

All medium to large-sized tech companies use performance management processes. The only places I'm aware of which operate without such processes are early-stage startups which are too young to need them, and small companies whose founders think that things run fine anyway, or who don't want to deal with this process.

3. Budgets

So, now we've decided to use a performance management process, can we jump straight into it? Not so fast. Performance management *always* gets tied to budgets. This has been the most important learning from my years as a manager. During the first few performance calibration sessions, I was unaware of the significance of budgets. However, here's the thing about them:

There's always a budget, even if you don't know what it is. Offers extended, raises and promotion raises don't happen in vacuum: they occur within the context of a budget. At more mature companies, budget planning typically takes place once or twice a year, involving the Chief Financial Officer (CFO), the CEO and others who need to have input. For a small startup, the budget could simply be in the founder's head or a small spreadsheet, upon which they tally up monthly and annual costs, and do the math on how much of a raise is affordable without eating too much into the remaining runway.

Budgets are better understood and enforced at companies with a CFO, whose job is to make the company's finances predictable. There's no better way to make costs predictable than to set a budget cap that can't be exceeded. Salaries, bonuses and staff pay rises are all costs which companies should want to budget for.

When it comes to budgets, we need to break them out into three categories:

- Cash bonus budgets
- Equity budgets
- Salary raise budgets

Cash bonus budgets

Cash bonuses are usually awarded once a year at most companies, together with salary raises, and after annual performance reviews. While most of Big Tech award one annual cash bonus – true for Google, Meta, Microsoft, Apple – there are other companies which might issue cash bonuses twice a year. Skyscanner did biannual bonuses, while some companies – usually outside of tech – may even issue quarterly cash bonuses.

A company which used to issue cash bonuses on a quarterly basis is Abnormal Security, which is an AI-based security email platform. The company was founded in 2018, and was last valued at \$4B in May 2022, when it raised more than \$200M of funding. The company issued cash bonuses every quarter, and advertised this unusually frequent bonus cadence to aid recruitment. However, this November, the company changed its bonus cadence to be annual, more in line with standard practice.

Bonus budgets can be set in a few ways:

1. Communicated upfront. Managers and employees are told what their target bonus for the coming year looks like. This is how Meta and Uber operate.

At Uber, every employee would be told the exact amount of their target bonus. This target bonus was what they expected to get if they performed to expectations. Although not communicated to all employees, there was an overall budget. In some years, this was 120% of the bonus budget and in other years it went higher, to around 130-135%.

At Meta, cash bonus targets are also communicated clearly, as detailed in [Inside Facebook's engineering culture](#). For example, IC5 (senior engineer) would have a 15% annual salary target, while an IC8's (roughly equivalent to principal engineer) would be 30%.

Meta goes further, and the multipliers are also [clearly communicated](#). For example, EE (exceeds expectations) is a 125% multiplier, while RE (redefines performance expectations) is a 300% one. Most companies don't communicate multipliers as clearly as Meta does, though, even within Big Tech.

2. Not communicated upfront. The majority of tech companies don't communicate cash bonus targets or budgets upfront. By not doing so, they have greater flexibility to set this budget, based on how the company is performing.

Even when a company does not communicate target bonuses, people build up a sense of what to expect, over time. Getting less cash bonus one year versus the former year will disappoint people.

The biggest upside of not communicating bonuses is that the company's leadership doesn't really need to plan ahead. A smaller downside is that when people don't know what size of bonus to expect, some might decide to not wait around to find out and choose to switch jobs, rather than wait months for bonuses to be awarded.

3. A hybrid: some is communicated. Some companies choose a middle path, of communicating some but not all details of what to expect.

For example, Uber could be seen as a bit of a hybrid. They do communicate the target bonus, but don't set expectations on how multipliers are used when it comes to being above or below expectations.

At Skyscanner, the targets were not communicated, but upon receiving cash bonuses, people were told what the multiplier was, between 1.0 and 1.5. This gives people a sense of how far above or below expectations they are. At the same time, the company has flexibility in setting the bonus budget, until the last minute.

Other businesses might communicate bonus targets upfront, but they also set a "company multiplier." For example, Meta also does this. Based on how well the company performs, the target can be adjusted up or down for everyone.

4. No cash bonus. Some companies simply don't pay cash bonuses. This might be a one-off: for example, during a bad year for the company. Or it can also be by design. Netflix is a better-known tech company which doesn't pay any form of cash or equity bonuses. In return, the company pays top of the market compensation, all of it in cash.

In The Scoop #29, we covered details of [Netflix's recent compensation philosophy changes](#) and how, because it paid all-cash, Netflix went from paying near the top of the market, to paying the very top of the market, due to its compensation structure containing the least equity among Big Tech.

Amazon is another company which does not issue cash bonuses, although it does issue equity refreshers over the years – but not for everyone.

Equity refresh budgets

Most of Big Tech and the majority of publicly traded tech companies allocate equity to all software engineers. The majority of such companies annually issue equity refreshers and the details vary by company. For example, Apple will issue an equity refresh and a cash bonus every year, to all software engineers of good standing. Meanwhile, Amazon will issue equity refreshers only to a subset of employees, based on their outstanding amount to vest and performance feedback.

The communication and budgeting of equity refresh budgets is done similarly to cash budgets.

1. Communicated upfront. For equity, it's exceedingly rare for refresh targets to be communicated. Meta is the only major company I'm aware of which communicates its equity refresh targets, which are [surprisingly generous, as I cover the specific numbers](#).

2. Not communicated upfront. The most common approach by most tech companies. Internally, there will be targets, usually visible at director, or VP-level. It's common enough that even line managers don't have access to this information.

3. No equity refresh. Many companies don't award equity refresh to the majority of staff. This is common for early-stage startups, where equity refresh is only given when people vest out their original package.

4. Only some people get equity refreshes. Some other companies outside Big Tech also limit the percentage of staff eligible for equity bonuses. For example, a director at Opentable told me a few years back that they were allowed to select 20% of their staff for an equity bonus, on top of the cash bonus.

Salary raise budgets

The one budget which is never communicated in advance, is the salary raise budget. This is because salary raises are linked to broad market conditions, like what competitors are doing, inflation, and the implications for retention.

Until 2022, in places with relatively low inflation such as the US or Europe, typical salary raises in tech averaged around 2-4%. In countries with already high inflation, like India or Argentina, the raises were much higher, of course. We covered [Microsoft's](#)

[compensation numbers](#) back in August, when this difference was clear: raises in India without promotions were at 13%, they were at 8% in the US, and in some European countries were as low as 5%.

Budgets for promotion raises are usually a separate category. Upon a promotion, a larger raise is expected. At most tech companies, a promoted person should expect at least an 8-10% raise. However, at company-level, there needs to be a budget for spending more on salaries, which will mean there will be a target for either the percentage of staff who can be promoted, or what amount can be spent on raises following promotions.

The amount to invest in raises is usually budgeted at the end of the preceding financial year. In mid-sized and large companies, the budget is usually set using market information for each role and job title. For example, a company will gather compensation data for software engineers and engineering managers, using a mix of:

- External salary data
- Obtaining a benchmark for competitors: e.g. those in the local or global markets, and based on a bucket of similar companies
- Data from candidates who rejected offers, and which offers they took
- The compensation for new hires

Smaller companies tend to rely on freely available data, such as public salary reports and sites which publish salary information.

Budgets will be the constraints which drive many calibration decisions. They especially influence any forced curve or bucketing, as we'll cover.


4. Performance calibrations

You're an engineering manager at a company with a structured performance process. As part of this process, managers do some preparation and employees might also prepare things like self-reviews, and could give peer reviews to colleagues. The manager goes through all the feedback and readies their assessment of their people.

However, how does this feedback turn into a cash bonus, equity bonus, or a salary raise? Enter the performance calibration.

The goal of the performance calibration is to ensure people across the organization who have performed similarly, get financial outcomes that reflect this as fairly and consistently as possible in the forms of a cash bonus, equity refresher, salary raise.

Performance calibrations: at a glance



The goal of performance calibrations is to identify groups of people based on performance and results, group people with comparable performance across the business, and commensurately reward similar performance.

Although not a goal, a side effect of a performance calibration is how it can help calibrate performance expectations across the group that does this session. Engineering managers will each interpret what a "high performer" and a "low

performer" is, somewhat differently. In a session, where they calibrate members of their respective teams, they'll likely come closer in the interpretation of expectations.

Depending on the company's stage, here's what calibrations may look like.

1. Early-stage: ad-hoc. At small or early-stage companies, it's typically the CEO or CTO getting together with a few key colleagues and going through the list of staff, which is how ratings are decided, and the sizes of bonus and pay rises. At small companies, there's usually not even buckets, as there's no need for them!

Usually, even the budget is up for negotiation, which again removes the need for planning too much for the calibration.

2. Mid-stage: engineering leadership running it. As the company grows, the formerly ad hoc performance calibration tends to be replaced by something with some formal rules in place.

It's common to introduce the concept of "buckets" of people across teams who performed similarly.

3. Late-stage: HR and CFO involvement. As companies mature more, HR typically gets involved in defining the process and ensuring its consistency. HR might start to create documentation and guidance around the process as well.

The CFO gets involved due to the setting of budgets, and how these constrain the amounts that can be allocated as rewards.

At later stage companies, promotions and the bonus + salary raise discussions are usually run as separate processes.


4. Company acquisition. When a company is acquired, the buyer will sometimes run a calibration session with the goal of leveling new staff, calibrating their pay bands and – occasionally – the "golden handcuffs" awards. This process usually differs for every acquisition.

5. Bucketing

Bucketing is a common technique when running calibrations at medium and above-sized companies.

Typical performance buckets within any calibration session tend to be:

- Below expectations. This bucket typically has a small number of people.
- Meets expectations. Usually the bucket containing the most people.
- Above expectations. A bucket with fewer people than in “meets expectations,” and often with more than in “below expectations.”



pragmaticengineer.com

A typical enough bucketing strategy, with three performance buckets

At many companies, there can be further separation of buckets, based on categories. It's common to have more granular categories for the “above expectations” bucket. Here's an example of a more granular bucketing system which I've observed. Note that such granular bucketing is rare:

- Below
 - Below most expectations
 - Below some expectations
- Meets

- Meets expectations
- New joiner – cannot rate
- Above
 - Exceeds some expectations
 - Exceeds

Facebook is a company which also has granular buckets. [We previously covered its categories](#), which are:

- Below
 - Does Not Meet Expectations
 - Meets Some Expectations
- Meets
 - Meets Most Expectations
 - Meets All Expectations
- Exceeds
 - Exceeds Expectations
 - Greatly Exceeds Expectations
 - Redefines Expectations

Performance buckets and promotion eligibility can be connected. At most companies, employees eligible for promotion are expected to have been in an “exceeds expectations” bucket during the last performance cycle, or are expected to be in there during the current one. Even at companies where it’s not a hard rule that those going for promotion should be in one of the “exceeds” buckets, it’s unlikely anyone gets promoted who is “only” in one of the “meets” buckets.

6. Stack ranking approaches

So, how does a large group of people get “bucketed” into the categories which the company defines? There are several approaches, though all of them are versions of

stack ranking. Like it or not, calibrations are about comparing people's past performance to their colleagues' and trying to ensure they are recognized fairly for it, in relation to their peers.

Forced stack ranking is an approach popularized originally by General Electric (GE) and was most notably used in the tech industry by Microsoft until 2013. Forced ranking means buckets have a fixed percentage to fill, which percentage is not negotiable. I worked at Microsoft in 2012, when there was this stack ranking. The buckets were:

- Low performer: 10%
- Middle: 70%
- High performers: 20%

There were a total of 5 levels within each bucket: 1 being the highest, and 5 the lowest. Needless to say, people wanted to avoid being in the bottom 10%, and hoped to end up in the top 20%. It was the same for me and I felt oddly disappointed to learn during my 2013 performance review that I was in the middle 70%.

There are two main benefits to forced ranking, which both reflect the company leadership's point of view:

1. Very easy to budget for. You can tell exactly how much you need to spend on various bonuses and raises, because the distribution is fixed.
2. Almost certainly guarantees the bottom 5-10% of staff leave, often voluntarily. Those who get the bottom rating are either people likely to be managed out, or who take the hint and seek new opportunities.

The downside of the forced ranking is that it can go wrong. As long as underperformers make up 10% of staff, and those people go into the appropriate bucket, then the system works well. However, if kept in place, one of two things will happen:

1. There are less than 10% of people who are underperforming. Because the quota must be met, some people who meet or even exceed expectations will be labeled as underperformers. Who's at risk of this? Well, typically it's those whose

managers aren't influential politically. Those people will then typically leave, even though they were not really underperformers!

2. Managers ensure there is always at least 10% of employees underperforming. How? By deliberately hiring people who can be "sacrificed" during calibration. Read more [about internal politics for software engineers and managers](#).

Stack ranking could have contributed to Microsoft's business stalling. An article in Vanity Fair in 2012, titled "Microsoft's lost decade," said:

"At the center of the cultural problems was a management system called "stack ranking." Every current and former Microsoft employee I interviewed—everyone—cited stack ranking as the most destructive process inside of Microsoft, something that drove out untold numbers of employees. (...)

In the end, the stack-ranking system crippled the ability to innovate at Microsoft, executives said. "I wanted to build a team of people who would work together and whose only focus would be on making great software," said Bill Hill, the former manager. "But you can't do that at Microsoft."


Coincidentally or not, after removing forced stack ranking in 2013, and appointing Satya Nadella as CEO in 2014, Microsoft broke out from the stagnation of 1999-2013, and resumed rapid growth.

Most tech companies avoid the rigid stack ranking which may partially explain Microsoft's failure to grow as a business. So, what do they use instead?

Ranking without specific *upfront* targets is a common approach. The buckets are defined, but managers are not pressured to meet specific targets for each bucket. At least, this is the idea.

In reality, there will always be some target, even if it's invisible to most managers. This is because there's a budget that drives this target. Still, the targets can be flexible or up for negotiation, even as the budget is fixed.

Stack ranking without upfront targets*



* save for the very bottom, and very top, usually

pragmaticengineer.com

Stack ranking without upfront targets. In this setup, there is usually still a target for the bottom and the top buckets, which may not be communicated upfront. But there is no such target for the buckets in between, other than that the distribution needs to fit the budget.

In most cases, the only buckets with a fixed target are the bottom and top ones. The bottom bucket typically has a 4-5% minimum target, and the top bucket a maximum of around 5%. Why set targets for these buckets? Because without targets, most managers would not nominate anyone on their team for the bottom bucket because this almost certainly causes attrition, and they would push as many of their people as possible into the top bucket.

This invisible target typically gets enforced only for larger groups, of around 50-100 people. At Big Tech, this is either at Director, Senior Director or VP-level, depending on the size of the population.

What about groups of less than 50-100 people, such as the staff engineer or director levels? For these levels, different and less strict guidance is often in place.

To be clear, there are companies where there really is no target. I've talked with senior managers and directors working in such environments, who shared that in the absence of targets, managers tend to push their direct reports a notch higher than where they should be. This is either because they want to avoid problematic situations in future, like questions about why an engineer got "meets" twice in a row, and whether they are performing. Or it's because they want to avoid a difficult conversation with an engineer about why they're "only" meeting expectations.

Ranking teams first, then assigning budgets. An interesting variation on stack ranking is to start by stack ranking teams based on delivered business results and key performance indicators (KPIs.) Teams are typically be bucketed into one of these categories:

1. Star (well above expectation)
2. At expectation
3. Below expectations

Depending which category a team is assigned, they get a different budget or distribution target to work with. "Star" teams get more budgets, or a distribution target where they can have more people in the "top performer" range than teams in lower buckets.

Assigning different budgets to teams based on their ranking does explicitly happen at some organizations, engineering directors and CTOs have told me. There are places where, although not explicit, it's implicit that some organizations which are profit centers, do get more budget to promote or reward top performers, compared to cost centers. We covered [*profit centers and cost centers in an earlier article*](#).

The positive side of ranking teams is that this approach can reward business outcomes better than a generic stack ranking. With a generic stack ranking, engineers getting the highest bonuses and pay rises may be those doing the most complex work, but not the work which moves the needle for the business. Ranking teams creates incentives for engineers to move to areas which are more important for the business, knowing it's easier to get rewarded there.

However, by allocating a different budget to each team, the downside is that good work won't be rewarded to the same degree across teams. Teams which are cost centers will likely always struggle to win higher allocations of budgets and bonuses.

How is going over budget handled? In the end, the calibration results – and bonus allocations or distribution of buckets – need to fit the budget allocated for the organizational unit. But what happens when this target becomes challenging to meet? Then exceptions get made, typically.

Organization leaders will attempt to negotiate with the budget owner on whether they can expand their budget. At smaller companies, this will mean talking with the CEO, and at larger companies it typically involves the VP-level. At Big Tech, budgets usually need to be adhered to at either the VP or director-level. However, a team or group can still get a larger budget at the expense of another getting less.

It is during these negotiations that the implicit "team impact ranking" can and does come into play. It's not uncommon for a VP to tell a cost center group that they need to reduce their top performer budget. What's unsaid is that this is because the budget of a profit center has increased.

Exceptions apply not just to bonus budgets, but also to promotion budgets. This is also why at large companies, managers are told to not communicate promotion outcomes until those promotions are reviewed and signed off by the C-level. Promotion budget adjustments can happen at the last minute, and I've also observed previously approved promotions reversed, due to budget constraints.

Takeaways

In this article, we've covered the systems that govern how performance review calibrations *need to work*.

Budgets are the biggest constraint for performance reviews or promotions. The money for bonuses, pay rises or promotions needs to come from somewhere, and someone must sign off the budget. Another incentive for owners of a company to put budgets in place, is to reduce spending on financial rewards.

Without any budget in place, engineering leaders shared that managers tend to inflate ratings and assign more generous rewards; in some cases, just to avoid awkward conversations. While this outcome is favorable for employees and managers, it is less beneficial for the business and its shareholders. Having a budget is a middle ground for offering financial rewards, but with sensible constraints upon their size.

All stack ranking approaches have trade-offs. All performance review processes end up comparing people, in some way. In the end, there will be a kind of ranking. This might be:

- Forced stack ranking, where people are ranked according to a fixed distribution. This makes budgeting easy, but most tech companies avoid this approach, after Microsoft's decade of stagnation was tied to its forced stack ranking approach – even though it's unclear how much to blame for sluggish business performance the forced stack ranking system was.
- Ranking without upfront targets is an approach many companies follow. There might be targets for the bottom and the top buckets, but the rest of the distribution is decided by the calibration group.
- Teams receiving different budgets based on their own "ranking," is what happens at some companies. Even at places which follow ranking without upfront targets, organizations that are profit centers might be able to lobby for budget exceptions.

There's no "perfect" choice for how to reward people with bonuses and pay rises. This means there's also no ideal solution for deciding how to set budgets, or what type of stack ranking approach to choose.

It is useful for you, as a manager, to understand which approach your company takes, as it has implications for the performance calibration process. If you're an individual contributor, it's educational to try and get information about how budgets work, and what happens during performance calibrations. Although you'll have no direct input into these processes, knowing what's rewarded and why it is, can help you develop an understanding of the system and make sense of the outcomes.

This was part one of a deep dive into performance calibrations at tech companies, in which we've covered the strengths and limitations of these calibrations. In part two,

we'll jump into what actually happens during a Big Tech calibration, and the strategies and politics which swirl inside the room when it takes place.

Thanks to [Andrea](#), [Jacek](#), [Jo](#), [Michał](#), [Karthik](#) and [Stefano](#) for their feedback and for reviewing an earlier draft of this article.



82 Likes

5 Comments



Write a comment...



Brendan Dec 6, 2022 ❤️ Liked by **Gergely Orosz**

> At Uber, every employee would be told the exact amount of their target bonus. This target bonus was what they expected to get if they performed to expectations. Although not communicated to all employees, there was an overall budget. In some years, this was 120% of the bonus budget and in other years it went higher, to around 130-135%.

Does this mean the target bonus for every employee exceeded the bonus budget by 20-35%?

❤️ LIKE (2) 💬 REPLY ...

3 replies by **Gergely Orosz** and others



Andrew Lee Dec 18, 2022

Great article — this has mirror my experience so far working in tech in California. I recently subscribed just to be able to read this. I found this through Kent Beck's substack to which I'm subscribed too.

Potentially a typo - while RE (redefines po;loexpectations) <-- the po;10 thing

❤️ LIKE 💬 REPLY ...

3 more comments...

© 2023 Gergely Orosz · [Privacy](#) · [Terms](#) · [Collection notice](#)
[Substack](#) is the home for great writing