THE PULSE

The Scoop #41: Meta's Big Restructuring to Come

Also: Amazon doubling down on RTO. And is ChatGPT causing panic within Big Tech?

GERGELY OROSZ

MAR 16, 2023 · PAID

70 8 1 Share

***The Scoop** is a bonus series covering insights, patterns, and trends I observe and hear about within Big Tech and at high growth startups. Have a scoop to share? Send me a message! I treat all such messages as anonymous.*

*The Scoop sometimes delivers first hand, original reporting. I'm adding an **'Exclusive'** label to news that features original reporting direct from my sources, as distinct from analysis, opinion, and reaction to events. Of course, I also analyze what's happening in the tech industry, citing other media sources and quoting them as I dive into trends I observe.*

Today's topics are:

1. **Meta: a second round of layoffs and restructuring to come.** The tech giant announced it will let another 11,000 staff go, and is also closing 5,000 open positions. I talked with software engineers there to get a sense of the mood inside the company, and what they expect will happen. ***Exclusive***.
2. **Amazon doubling down on RTO (return to office).** The tech giant is not blinking: after the unprecedented pushback on its 3-day return to office policy last month, Amazon just made it clear that the policy is going ahead. I talked with engineers inside the company for their take**. *Exclusive.***
3. **Is ChatGPT causing panic within Big Tech?** Large-scale language model ChatGPT is taking the industry by storm, despite Big Tech companies having

invested heavily for years in natural language processing (NLP) teams. Are these teams now panicking, and can we expect Big Tech to increase or decrease investment in NLP? *Analysis*.

4. **How much do full-remote take-home exercises pay?** During their interview process, several full-remote companies ask software engineers to complete take-home exercises, which can take several days to complete. Among companies which pay candidates for their time, what are the rates? I gathered details from engineers. ***Exclusive.***

1. Meta: a second round of layoffs and restructuring to come.

On Tuesday 14 March, Mark Zuckerberg sent an email to employees titled "Update on Meta's Year of Efficiency." Three paragraphs into this email, he dropped a bomb:

> "Overall, we expect to reduce our team size by around 10,000 people and to close around 5,000 additional open roles that we haven't yet hired."

Layoff notifications for engineering are to follow in late April, while non-tech teams will get notifications in late May.

Until last year, Meta had not done a single layoff in its 18 year-long existence, and was optimistic about avoiding job cuts, despite its revenue slowdown. However, on 9 November 2022, the company eliminated 11,000 positions: about 10% of software engineers were let go, with recruitment hit much harder, with about 50% impacted.

Back then, I concluded:

> "What does this layoff at Meta mean for the rest of the industry? Unfortunately, nothing good. Meta was the holdout company whose founder CEO could afford to not do layoffs, despite mounting investor pressure. The rest of Big Tech does not have this luxury, and I expect these cuts now erase the stigma of layoffs within Big Tech. We can, sadly, expect that in the coming year, more Big Techs will reach to terminate people as a quick and convenient way to reduce costs and increase profits."

When I wrote I expected more Big Techs to do redundancies for cost savings, as Amazon and Google have, I didn't expect it to be Meta less than 6 months later.

The conventional wisdom on layoffs is to cut once, cut deep, and cut fast. This is because layoffs are devastating to morale and the only thing worse than one layoff round, is two of them. With the first round of layoffs, there's hope that the "remedy" will work, however bad it is. A second set of layoffs indicates that leadership failed with the first round, so who knows how many more will follow?

To Mark Zuckerberg's credit, he hasn't not skirted about the reason for this second layoff round in his email to employees. This is because these cuts are the start of a corporate restructure which started when Meta decided to reduce its number of middle managers a month ago.

In the latest email, Zuckerberg outlines how he plans to change the company:

- **Flatter is faster.** "In our Year of Efficiency, we will make our organization flatter by removing multiple layers of management. As part of this, we will ask many managers to become individual contributors. We'll also have individual contributors report into almost every level — not just the bottom."
- **Leaner is better.** "A leaner org will execute its highest priorities faster. People will be more productive, and their work will be more fun and fulfilling. We will become an even greater magnet for the most talented people. That's why in our Year of Efficiency, we are focused on canceling projects that are duplicative or lower priority and making every organization as lean as possible."
- **Keep technology the main thing.** "As we've grown, we've hired many leading experts in areas outside engineering. This helps us build better products, but (...) we're focusing on returning to a more optimal ratio of engineers to other roles."
- **Invest in tools to be more efficient.** "[We will be] investing in tools that will make us most effective over many years, not just this year — whether that's building AI tools to help engineers write better code faster, enabling us to automate workloads over time, or identifying obsolete processes that we can phase out."
- **In-person time helps build relationships and get more done.** "Our early analysis of performance data suggests that engineers who either joined Meta in-

person and then transferred to remote or remained in-person, performed better on average than people who joined remotely. This analysis also shows that engineers earlier in their career perform better on average when they work in-person with teammates at least three days a week."

Having read the full email, here is my prediction about what will happen at Meta:

- **A purge of middle management, mostly outside of engineering.** At the end of 2022, Meta had around 87,000 employees, and based on insider data I obtained, the average "distance" from an individual contributor to Mark Zuckerberg was a whopping 9 levels. In comparison, many teams at Apple have around 6 levels between engineers and Tim Cook. With Meta eliminating about 20,000 roles between the November layoffs and these upcoming ones, the workforce will likely be nearer 70,000. I'd be surprised if the average reporting distance was wider than 7 levels, and maybe closer to 6, when this is done. As noted in the article, [Meta to have fewer managers: insider comments], engineering is expected to be hit less by a middle management cull, as engineering managers tend to have more direct reports than managers in other domains.
- **Non-engineering will be hit much harder by this round of layoffs.** Zuckerberg was pretty clear in his assessment: Meta has hired too many people outside of engineering, distracting from its core values. In the early days of Facebook, engineering was at the center of everything, and these cuts surely indicate some return to this. Also, Zuck makes clear engineering tooling investments are here to stay.
- **A return to the office for 3 days/week is around the corner.** Meta was the last Big Tech company to have not mandated a return to the office. Apple, Google, and Microsoft have all adopted a 3-days-per-week policy, as has Amazon despite [an unprecedented pushback]. This email sets out the core argument why Meta will require, at the very least, less experienced engineers to be in the office 3 times per week. I expect this policy to eventually be enforced for more junior levels with a more flexible approach for more senior folks.
- **Meta's business is still at the advertising market's mercy.** 97% of revenue comes from ads, which contributes to [its historic growth challenge]. It's still unclear to me how Zuckerberg plans to change this, or if he really has a plan. The

> Metaverse investment continues to be a long-term bet that Meta will become a hardware platform, but it's hard to see this investment gaining mainstream traction. In this email, Zuckerberg shifted his tone and wrote how "our single largest investment is in advancing AI and building it into every one of our products," which suggests the goal is to make products like Instagram, Facebook, and Messenger even stickier for users, in order to attract more ad spend. But is there any plan to diversify beyond ads? If there is one, it's not clear.

I talked with current software engineers to discover what the mood is like inside Meta. Here's what I gathered from five:

The layoff announcement was not a big surprise for most engineers. This was because on Blind – an anonymous internal network – people got the warning ahead of time in the internal Meta lounge. A user called "bozztank" had shared the news of the November layoffs on Blind a week before they happened. And bozztank posted again a week before Zuch sent out his latest email, telling people that layoffs are coming. Although Blind is an anonymous network where users can change their names, a few accounts tend to build up a reputation for sharing reliable insider information, of which bozztank is now one.

All engineers I talked with appreciated the heads-up from the CEO that layoffs will happen. Conventional wisdom around layoffs is that "cutting fast" is usually a better approach because once a layoff is announced, work grinds to a halt. Perhaps this is a reaction to having gone through surprise job cuts quite recently.

Morale is predictably down. But to me, it feels different to the low morale around last November's layoffs. Here's how a software engineer put it:

> "Everyone in my team seems pretty 'off'. I would say morale 's better than in November, but not significantly so. But now that we know what to expect, I think it'll be business as usual until April.
>
> My team was pretty upset after November. We had a handful of people let go during those layoffs. For a few weeks, nobody on the team did anything meaningful, work-wise."

The consensus among most engineers is that as long as you are on a reasonably important topic and "meet expectations," you should be fine. Talking with someone in engineering leadership, they shared that there are not many engineering managers marked out for redundancy. There's much more "flattening" happening: like M1s to report to an M2, instead of to a fellow M1.

Engineers below expectations and getting a MM, MS or DNE, are worried about being let go. The same is true for engineers working in cost centers, those projects not generating revenue. Engineers outside of Metaverse projects speculate that a big cut could be made to the Metaverse, given Zuckerberg has made clear that AI is Meta's main priority, not VR. Wondering whether or not an engineering team is a profit center or not, is a new problem. Until now, all of engineering was considered a profit center. This new round of cuts should make the distinction clearer: projects which Meta's leadership consider as cost centers are likely to see heavier cuts.

I sense a major shift in sentiment: more software engineers no longer feel Meta is 'special' or 'different.' Before the first round of layoffs, there was faith that Zuck saw things differently from every other tech leader. However, I now sense resignation: people accept Meta is no different to Microsoft, or even IBM, in how the business works. Here's how engineers described it:

> "Now Zuck is just like any other CEO in how he runs the company. It is disappointing."
>
> "I don't care anymore what happens. If they lay me off, it's the company's loss lol."
>
> "The situation is a mess. I don't even think layoffs are solving the core problem we have."

The engineers who feel most bitter are all recent joiners. I talked with two engineers who joined within the past nine months. Both left behind troubled scaleups which they expected to struggle, and expected Meta to be a safe haven during turbulent times. What they did not expect was for the storm to be at Meta, too. One of the engineers put it like this:

> "Bootcampers are blocked from joining teams until after the layoffs. We delayed starting dates for new joiners by 4 months, and it still wasn't enough to allow them to join dates. I don't understand why those bootcamper offers were not just rescinded: people are left hanging in limbo, doing nothing but waiting.
>
> It's disappointing to see what will account for a 25% layoff for Meta - 13% in November, and now again that much. Keep in mind, the company is still wildly profitable and just signed off on $40B in stock buybacks."

In many engineering teams, morale seems to be better this time, than it was last November. I don't know if it's clearer messaging about engineering teams being safer, or just team dynamics. But in groups on Messenger and WhatsApp, I talked with engineers who sound less concerned.

Meta's drastic cut is not great news for the tech industry, and other companies could follow suit. Mark Zuckerberg seems determined to pull back Meta to become not just a leaner organization, but an engineering-led one. The reason this news isn't great for workers is that there's a fair chance other companies will follow. After all, Twitter let go 70-80% of staff in order to cut costs and it seems this inspired other businesses. And although Twitter's move was extreme and not one I'd advocate, the cuts at Meta feel… surprisingly sensible? Zuckerberg is making it clear that the biggest problems are:

1. Excessively long reporting lines
2. The explosion of "busywork" that comes with low-priority projects and their overhead
3. No longer being a tech-first company

Zuckerberg's focus on returning Meta to its tech first-roots is likely a good thing for engineers, but less good for non-tech roles. The focus on operating in a leaner way will also be good for engineers; those who stay, that is. And a leaner organization will attract more people. As Zuckerberg said:

"A leaner org will execute its highest priorities faster. People will be more productive, and their work will be more fun and fulfilling. We will become an even greater magnet for the most talented people. That's why in our Year of Efficiency, we are focused on canceling projects that are duplicative or lower priority and making every organization as lean as possible."

2. Amazon doubling down on RTO (return to office)

One month ago - on 17 February - Amazon announced its new remote work policy. From 1 May, employees are expected to spend at least three days per week in the office.

I shared insider details on the unprecedented pushback that this new policy has received. A big part of the pushback came because of three big reasons:

1. Many employees accepted offers from Amazon with the understanding that their position will be full remote. About 30-40% of Amazon's current staff was hired in the past 18 months, based on an engineer I talked with, looking at the "Old Fart" ranking - a system that shows what percentage of employees were hired after you.
2. Plenty of employees are not within commuting distance from their assigned office.
3. From 2021, directors had the autonomy to decide if their teams worked as remotely, hybrid, or in-office. When hiring employees, directors often assumed that this setup would stay permanently.

At the time of publishing, more than 30,000 employees signed a petition asking for this new RTO policy be reversed, and more than 32,000 people joined an internal Slack channel to protest against the RTO policy change, I confirmed with a current software engineer. As we covered in The Scoop #39, this petition was addressed to the S-Team and reads:

> "We, the undersigned Amazonians, petition to reverse the RTO initiative announced Friday, February 17, 2023, and propose Amazon reinstate the remote-work policy that empowers teams and individuals to determine and adopt the work style that works best for them - remote, hybrid or in-office."

Engineers I talked with were hopeful that Amazon's leadership team - the S-Team - would reverse the policy or at least soften it. Until now, the S-Team did not communicate further about these changes.

Until today, that is. Today - on 16 March - Amazon updated their FAQs about RTO. There is no softening of the policy: Amazon is doubling down on it. The company confirmed that building assignments for employees will be published by 14 April - referring to the building those employees need to commute to. They also published answers to 28 frequent questions, including these (emphasis mine):

> Q: "If I now live far away from my assigned office, do I need to come into the office 3 days a week once my building is ready?"
>
> A: "Yes. **We expect all employees to return to their assigned office at least three days a week** when their building is ready. If there is a reason in which you cannot do so, please discuss with your manager to explore options such as allowing extra time until you can relocate back to your assigned area of transferring to another team near your location."

The question about when your team is remote but you are not:

> Q: "What if I'm the only or one of just a few people from my team in the office I'm assigned to - do I still need to go in at least three days a week?"
>
> A: "Yes. We believe that employees are much more likely to understand our unique culture and become part of it if they are surrounded by other Amazonians in person, even if not on their immediate working team. **We encourage managers to work towards having as many of their team members together in one physical location as possible**. We recognize that this will take time, and in some cases, managers may choose to make an exception and designate certain individuals as remote, but this will be a very small minority of the company overall."

Sales and Customer Support will probably be exempt from the RTO requirement: however, employees do not yet have details on exactly what teams. For people outside Sales and Customer Support - such as Engineering - exceptions need to be requested to work remotely. If the exception is not related to a disability accommodation: this exception requires S-Team approval - which, practically, means an SVP or VP-level approval.

On exceptions:

> Q: "Are there exceptions to returning to the office three days a week? If I'm unable or unwilling to return to the office, how do I request an exception to work remotely?"
>
> A: "Yes. As there were before the pandemic, there will be exceptions for certain roles and functions that require remote status. For example, some in sales or customer support have historically had remote-working status and that will stay the same moving forward. If your role or function does not allow for remote-work status, you can request an exception to work remotely by starting with your manager to discuss your situation (...) An S-team member approval will be required for remote work status."

Talking with engineering managers at Amazon (called SDMs here), those SDMs did not have details on how exceptions work, and neither did their directors. Engineering managers I talked with hope to get more clarification on the exceptions process later on.

So what is the mood like inside Amazon? Grim, talking to engineers. On Slack, people are talking about resigning, and the topic of unionizing is coming up more frequently. Amazon is known for being a company that wants to stop unionization by every means possible - so bringing up unionization - or, as engineers, I talked to called it: 'the forbidden U word' is an unusually strong response.

The same engineers upset about the RTO policy here to stay are also talking about canceling their Prime memberships or boycotting Amazon.

What I hear from all of this is that many engineers find their situation hopeless: I am talking about people living too far from an office to make their commute practical. These people have no other option than to apply for a remote work exception - and hope they get it from the S-Team. Failing that: it will be a very long commute from May or having to find another job and resign.

What does Amazon's approach signal for other companies? I consistently pay attention to what Amazon does, as the company has read the market well for the past decades: both from a business perspective, and also from how to attract and retain software engineers - doing this in a frugal way.

The fact that Amazon's S-Team did not blink, and refuses to soften the return to work policy signals that they either don't expect much attrition, or have calculated with additional attrition. Given that the company is in the middle of cutting 18,000 positions, there's a fair chance that this RTO policy is a way to further increase attrition.

At the same time, looking at the market, it's hard not to ask the question: if people want to leave, where will they leave to? Meta just announced more cuts - and closing all 5,000 open positions. Dozens of tech companies have put similar, 3-day return to office policies in place starting from December, when an RTO wave started, as we covered in The Scoop #34.

There are, of course, still companies out there that are hiring for either full-remote or remote-friendly roles: a large number of startups and scaleups, and some better-known tech companies like Shopify, Spotify, GitLab or Automatic.

Still, I cannot help but observe that we have passed "peak remote work," and workplaces that have been office-first before the pandemic: most of them are reverting to a hybrid setup of 3 days per week in the office.

Just like Meta has hinted of doing exactly this, likely preparing for a policy change, also to target three days per week in the office.

3. Is ChatGPT causing panic within Big Tech?

ChatGPT and OpenAI have been dominating news about AI and language processing since the public release of ChatGPT in November 2022. In just 3 months, ChatGPT achieved 100 million active users, making it the fastest-growing consumer "app" in history, ("app" is apostrophized because ChatGPT is, really, a platform.) Anyone can try out ChatGPT, and developers can use ChatGPT via APIs. We've seen businesses adopting ChatGPT rapidly, including:

- The new Bing launched with ChatGPT in-built, in early February.

- Snap launched My AI, an experimental chatbot and Shopify launched Shop - a shopping assistant using ChatGPT, in late February.
- Duolingo introduced Duolingo Max, a learning experience powered by GTP-4, on 14 March.
- Instacart plans to launch "Ask Instacart" for customers to discover ideas for open-ended shopping goals, later this year.

My observation is that the developer community, as a whole, hasn't been this excited about a new technology in a long time:

Gergely Orosz
@GergelyOrosz

The way OpenAI is building excitement across devs is something I've not seen in a long time. A lot of this is deliberate, eg:

- The product launches with GPT-4 on the day of announcements
- The (not advertised) "exclusive" events w select CTOs being invited last minute to fly in

1:39 PM · Mar 15, 2023

841 Likes **37** Retweets

I last saw such widespread excitement in the 2010s for Apple's product launches; when mobile developers knew something big was coming out, and started using the new APIs announced at the WWDC, on the same day. Earlier, in the 2000s, Microsoft's announcements caused excitement among the Microsoft development crowd, to the point that several publications started specializing in Microsoft-only dev news.

Back to today, and most Big Tech companies have been investing in NLP for years. NLP stands for Natural Language Processing. In the book Practical Natural Language Processing by Sowmya Vajjala, Bodhisattwa Majumder, Anuj Gupta and Harshit Surana, the authors summarize the tasks and applications of NLP:

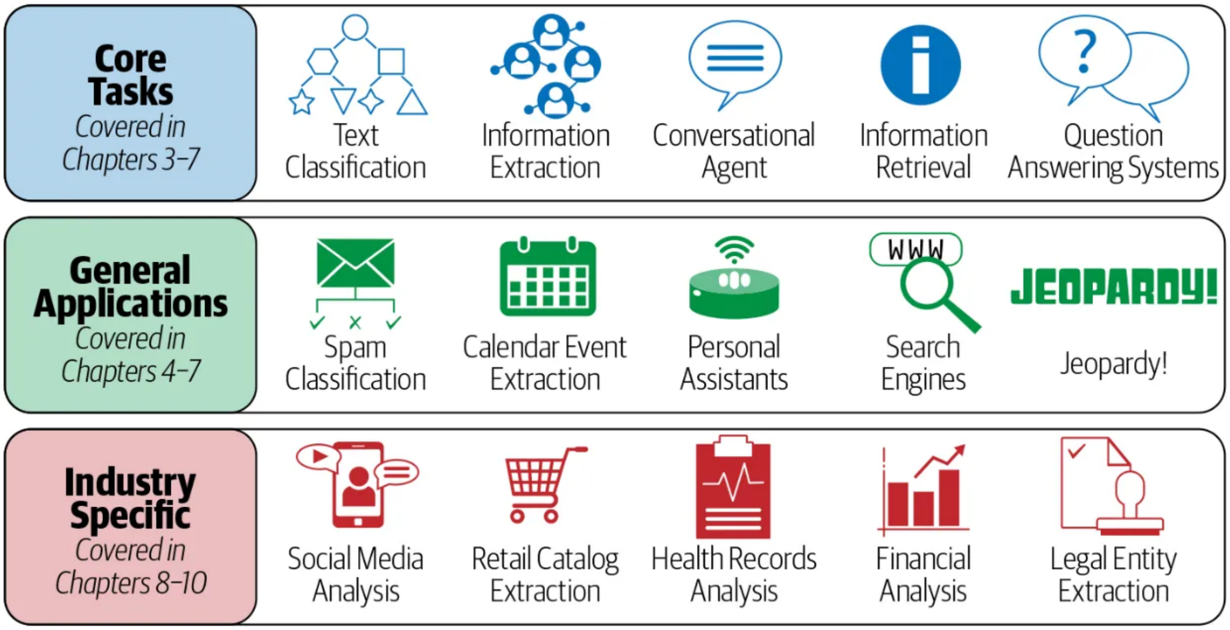

Figure 1-1. NLP tasks and applications

Categorizing NLP tasks and applications. *Source: Practical Natural Language Processing*

The book categorizes the most common approaches to NLP like this:

- Heuristics-based NLP: previous attempts, building rules for tasks at hand.
- Machine learning approaches like Naive Bayes, Support vector machine and others.
- Deep learning: recurrent neural networks, long short-term memory and others.

NLP use cases are practical for tech companies to improve products, from building conversational agents to extracting events from calendars. It's no wonder most Big Tech companies have been employing large teams to productionize NLP approaches to build into their products in the form of better search recommendations, better assistants, better analysis tools, etc.

What's different about ChatGPT, which is based on large language models (LLMs), is that it seems to have achieved the breakthrough of doing most NLP tasks as well as – and often much better – than the existing approaches. This looks set to accelerate with the release of GPT-4.

"Anyone else witnessing a panic inside NLP organizations of Big Tech companies?" This is the title of a new thread on Reddit, to which I was pointed by a Big Tech software engineer, who says this is what they're seeing at their workplace. The author of the post writes:

> "I'm in a Big Tech company working alongside a science team for a product you've all probably used. We have these year-long initiatives to productionalize 'state of the art NLP models' that are now completely obsolete in the face of GPT-4. I think at first the science orgs were quiet/in denial. But now it's very obvious we are basically working on worthless technology. And by 'we", I mean a large organization with scores of teams.
>
> Anyone else seeing this? What is the long-term effect on science careers that get disrupted like this?"

There are hundreds of responses in the thread from people working in the field. Here's a few I find interesting:

A current science intern and an ML engineer discuss how ChatGPT simplifies some common tasks:

> Science intern: "[I am] just an intern, but my previous work on NLP sentiment analysis is 100% obsolete. My previous seniors are very worried that most of their work is now just implementing GPT APIs.
>
> ML engineer: "What I found to work well:
>
> 1. Text extremely hard to classify -> GPT3.5 API with a prompt that explains the task in detail.
>
> 2. Text not THAT hard to classify + ability to deploy sufficiently large transformer -> train your own model.
>
> Lacking data? GPT3.5 till you collect enough training data."

On the importance of data versus models:

> "I see the same. The panic is entirely predictable. These models are seen as an existential threat to the things they have been working on the last few years.
>
> But there's a reason that productionization of those NLP models has taken years (quality, cost, and latency) and I don't think [ChatGPT] is magically going to fix those problems. In many cases [ChatGPT] is going to be worse than an in-house model because you can't control cost or latency, address quality issues, and now you have a dependency on an external API.
>
> [ChatGPT] will be disruptive in some areas but I don't think anyone really knows which use cases will become billion dollar profit machines and which ones are just money pit, tech-fever dreams.
>
> I think data is way more important than models and OpenAI has been very smart about data collection. I see no reason others can't catch up."

There could be an impact on academia, believes one commenter:

> "Panic inside NLP orgs of big tech companies? What about the panic at NLP departments in universities? I have witnessed my friends putting their work on PhDs go into despair after ChatGPT and now GPT-4. Quite literally, the majority of the research topics in NLP are slowly becoming obsolete in front of our eyes."

On the rest of the industry catching up with OpenAI:

> "I think it will be a game of catch up between the open-source models and OpenAI's models with the open-source models lagging 6-8 months behind."

The breakneck speed of applied innovation by ChatGPT will have an industry-wide impact. OpenAI has made large language models accessible to every software engineer, usable either through a REST API or through a chat interface.

When I first tried out ChatGPT, I had a "wow" moment which reminded me of when I used a touchscreen phone for the first time, or when I first played with the Ultraleap. It's not just me, but most engineers I know. Here's Dan Abramov, creator of Redux -and someone who was in the React.js documentary - on his ChatGTP experience:

дэн
@dan_abramov

chatgpt is completely blowing my mind rn

5:51 AM · Mar 16, 2023

598 Likes **19** Retweets

A software engineer at Meta whom I've been talking with, has wanted to get into machine learning for a good time, but it wasn't possible to transfer internally to an ML team since they don't have a PhD on this topic. He welcomes the change that's unfolding, saying:

> "I've been on a 24x7 LLM immersion for the past month. It feels to me that OpenAI turned NLP and ML development into something like React.js development.
>
> While I can see the discomfort about all the new people coming into the field from the established people: I am happy. These people were so gatekeepy: unless you had PhD papers you did not even get to work on, or be let close to deploy ML systems. OpenAI has turned the "open" part into a joke by not opening up their research. However, they *are* empowering devs by turning AI into 3 lines of JavaScript."

ChatGPT feels like it's bringing LLMs to engineers. Thanks to such rapid adoption and a very clear product-market-fit, I expect a response from Big Tech. Possibilities include:

1. Redirecting NLP efforts to LLMs, following the path that OpenAI has proved to be viable. Most large companies will not want to rely on an external vendor like OpenAI when it comes to technology that can be a core advantage. Google and Meta are already known to be working on their own LLM models, and it's safe to assume companies like Amazon and Apple are doing the same. Microsoft, of course, practically owns OpenAI, thanks to an exclusive long-term partnership. We only have rumors about this deal, which suggest Microsoft could own a 49% ownership stake in the company.

2. Reducing funding on NLP efforts unlikely to bring practical results which can compete with LLM outputs. This would be a natural reaction to a technological breakthrough: why keep investing in approaches that look unlikely to become profit centers, in light of innovations?
3. Smaller companies without the budget to invest in LLMs could decide to either license solutions like ChatGPT, or experiment with using or investing in open source approaches, such as Together.

AI and machine learning will keep being investment areas within Big Tech. Even in the midst of a Big Tech Hiring slowdown, these companies have made it clear they are actively hiring for AI-related areas. In October I wrote:

> "[Even during their hiring freeze] Meta will keep hiring for AI and machine learning (ML) roles, until it hits its targets.
>
> Google is doubling down in AI. Answering a question about how Google will respond to economic forecasts which suggest a recession is looming, the CEO responded by highlighting AI's role.
>
> Microsoft spelled out that it expects to invest more in data & AI (OpenAI relationship)"

This sentiment has not changed since. If anything, all three companies - Meta, Microsoft, and Google - seem to be investing in AI even more.

4. How much do full-remote take-home exercises pay?

I've observed several startups hiring for full-remote roles ask candidates to complete a take-home exercise. This exercise can last a few hours, to a few "trial days," and even a "trial week." One employer that limits the length of its exercises to about two hours is Sourcegraph, as we discuss this week in Inside Sourcegraph's Engineering Culture. However, when companies ask candidates to put in more than a few hours, most offer payment in return.

But what are the rates? I asked around and found two distinctive types:

1. Paying roughly the rate for contracting:

- $125/hour ($1,000) for a day-long exercise at a Seed-stage startup
- $100/hour ($1,600 for 2 days) at a Series A company
- $600/day (or $3,000 for the week) for a week-long trial week at a Series B startup
- The "normal salary" paid for a week. At a Series A company, candidates on the 1-week trial were paid the same as the job itself does. If it took longer to complete the work – for example, due to doing it part-time while working a current job – candidates were still paid the 1-week "earnings"

2. Pay around $20-40/hour:

- $25-40/hour at a well-known full-remote company with 500+ employees. Candidates can take as long as they need: some take a week, others up to 3 months. Candidates send a bill based on their time investment.
- $20/hour at an outsourcing company.
- ~$20/hr ($200 in Amazon gift cards for about 10 hours of work) for candidates who finished assigned tickets at a Series A company.

Balancing how long an interview should take is always tricky. In cases where both the candidate and the company invest equal time in interviews with in-person interviews, it's clear both parties devote time, energy and money to the process. However, when it's the candidate spending the most time and the employer very little – and mainly on evaluating results – then the balance is asymmetric.

I believe employers should compensate candidates for asking for more than a few hours of asynchronous work. Of course, this is if the company is serious about hiring strong software engineers. While some engineers with other options will likely pass on a lengthy take-home exercise, paying a contractor rate makes the exchange fairer, and the candidate less likely to drop out.

Also, paying a contractor rate means the company hiring about who does the take-home exercise. A startup CEO I talked with said their goal is to get to the point where 70-80% of these "final round" candidates get an offer. Right now, this company is at around 30-40%, and they want to improve: after all, it's costing them about $100/hour for not doing screening more efficiently earlier in the process!

Takeaways

There's never a boring day in tech - much less a boring week! This week, the two biggest pieces of news have been the collapse of Silicon Valley Bank and Meta's second round of job cuts, and the restructuring to follow. Perhaps a quiet news day would be welcome, right now.

I expect to see a renewed focus on efficiency and a "wartime" approach at many tech companies.

And looking at what's happening in the AI field, who can blame Big Tech for wanting to be leaner? OpenAI, which has less than 500 employees, is upending the industry and seemingly running rings around sluggish-looking Big Tech companies which have far bigger teams working on AI and NLP. It's another case of bigger workforces not necessarily equalling greater output. Big Tech is likely to respond.

Do you have tips, scoop to share or feedback on this issue? Send me a message. I treat all sources as anonymous.

Updates on 18 March: added details on building assignments communicated by 14 April, and more details on exceptions.

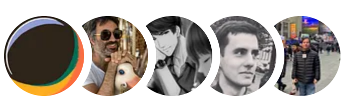

70 Likes · 1 Restack

8 Comments

Write a comment...

William Pittman Writes The Virtual Mentor Mar 16 Liked by Gergely Orosz

Re: ChatGPT

"My observation is that the developer community, as a whole, hasn't been this excited about a new technology in a long time"

Strong agree. I attended an in-person tech meetup in Seattle last week. Everyone I spoke with wanted to talk about ChatGPT. Every. Single. Person.

LIKE (4) REPLY

Jaime Buelta Mar 16 · *edited Mar 16* Liked by Gergely Orosz

Re: Amazon's RTO

This is ridiculous... Are they expecting someone that is in a different country than their team to go to the office to work? Just to be surrounded by other teams? I know someone in that situation, and it will be extremely awkward just to go to sit on a corner to take Zoom calls after two years working remotely

LIKE (3) REPLY

2 replies by Gergely Orosz and others

6 more comments...